

Degeneracy of the Inverse Problem: Uncountable Solutions from Dynamics

1 The Forward Model

The flyvis graded-voltage ODE for each postsynaptic neuron i reads:

$$\tau_i \frac{dv_i(t)}{dt} = -v_i(t) + V_i^{\text{rest}} + \sum_{j \in \mathcal{N}_i} \mathbf{W}_{ij} \text{ReLU}(v_j(t)) + I_i(t) \quad (1)$$

where τ_i and V_i^{rest} are cell-type parameters, $\mathbf{W}_{ij}$ is the connectome-constrained synaptic weight from neuron j to i , and $I_i(t)$ is the external stimulus.

The inverse problem: given observed trajectories $\{v_i(t)\}_{t=0}^T$, stimuli $\{I_i(t)\}$, and the binary adjacency matrix $\mathbf{A}$ (which edges exist), recover the connectivity weights $\mathbf{W}$, time constants τ , and resting potentials V^{rest} .

2 Where the Degeneracy is Hard-Coded

2.1 Step 1: Rearrange to a linear system per neuron

Rearranging Eq. 1 to isolate the incoming synaptic weights, we move them to the left-hand side:

$$\sum_{j \in \mathcal{N}_i} \mathbf{W}_{ij} \text{ReLU}(v_j(t)) = \tau_i \frac{dv_i(t)}{dt} + v_i(t) - V_i^{\text{rest}} - I_i(t) \quad (2)$$

Define the presynaptic activity $h_j(t) = \text{ReLU}(v_j(t))$ and the residual $b_i(t) = \tau_i \frac{dv_i(t)}{dt} + v_i(t) - V_i^{\text{rest}} - I_i(t)$. The constraint is then:

$$\sum_{j \in \mathcal{N}_i} \mathbf{W}_{ij} h_j(t) = b_i(t) \quad \forall t \in \{1, \dots, T\} \quad (3)$$

Stack these T timepoint constraints into a linear system: Let $\mathbf{H}_i \in \mathbb{R}^{T \times d_i}$ be the activity matrix (rows = time, columns = the d_i presynaptic neurons), $\mathbf{w}_i \in \mathbb{R}^{d_i}$ the incoming weight vector, and $\mathbf{b}_i \in \mathbb{R}^T$ the stacked residuals. Then:

$$\mathbf{H}_i \mathbf{w}_i = \mathbf{b}_i \quad (4)$$

This equation holds independently for each postsynaptic neuron $i = 1, \dots, N_{\text{post}}$, where $N_{\text{post}} = 13,697$ and d_i ranges from 1 to 208 (the in-degree distribution of the flyvis connectome).

2.2 Step 2: The null space

Any $\boldsymbol{\delta} \in \ker(\mathbf{H}_i)$ satisfies $\mathbf{H}_i \boldsymbol{\delta} = \mathbf{0}$. Therefore the perturbed weight vector $\mathbf{w}_i + \boldsymbol{\delta}$ produces exactly the same dynamics as $\mathbf{w}_i$:

$$\mathbf{H}_i(\mathbf{w}_i + \boldsymbol{\delta}) = \mathbf{H}_i \mathbf{w}_i + \underbrace{\mathbf{H}_i \boldsymbol{\delta}}_{=\mathbf{0}} = \mathbf{b}_i \quad (5)$$

So any perturbation in the kernel of $\mathbf{H}_i$ is invisible to the dynamics. How many such directions exist? The matrix $\mathbf{H}_i$ has d_i columns (one per presynaptic neuron, $d_i \in [1, 208]$). By the rank-nullity theorem (for any matrix with n columns: $\text{rank} + \text{nullity} = n$), these d_i columns split into two groups:

$$d_i = \underbrace{\text{rank}(\mathbf{H}_i)}_{\text{identifiable weights}} + \underbrace{\text{dim ker}(\mathbf{H}_i)}_{\text{free (null) weights}} \quad (6)$$

Global bound from SVD. An SVD of the full population activity reveals that the entire circuit’s activity lives in a low-dimensional subspace. The computation subsamples to 982 neurons (every 14th of 13,741) and 8,000 timesteps (every 8th of 64,000), then computes the full SVD (all 982 singular values):

Metric	Value
Neurons (subsampled)	982 / 13,741
Timesteps (subsampled)	8,000 / 64,000
SVD	full (982 singular values)
Effective rank at 90% variance	1
Effective rank at 99% variance	45

Since each neuron i ’s presynaptic activity $\mathbf{H}_i$ is a submatrix of the global activity, its rank is upper-bounded by the global effective rank: $\text{rank}_{\text{eff}}(\mathbf{H}_i) \leq 45$. This is an *approximate* bound—the 99% variance threshold is a convention, and the true mathematical rank (number of nonzero singular values) may be higher. Replacing exact rank with effective rank gives an approximate null space:

$$\text{dim ker}_{\text{eff}}(\mathbf{H}_i) \approx \max(0, d_i - r) \quad (7)$$

where r is $\text{rank}_{\text{eff}}(\mathbf{H}_i)$, the effective activity rank. This approximation predicts that perturbations along the neglected singular directions have negligible effect on dynamics. Summing over all 13,697 postsynaptic neurons gives a global estimate:

$$\text{total null dim} = \sum_{i=1}^{N_{\text{post}}} \max(0, d_i - r) \quad (8)$$

Sensitivity to the variance threshold. Since the 99% threshold is arbitrarily chosen, we report the null space estimate across several thresholds:

Variance threshold	Effective rank r	Null space dim	% identifiable
90%	1	420,415	3.2%
95%	2	407,427	6.1%
99%	45	115,223	73.5%
99.5%	90	26,413	93.9%
99.9%	238	0	100.0%

At 99.9% the rank (238) exceeds the maximum in-degree (208), so the null space vanishes by this metric alone. The 99% threshold ($r = 45$) gives the most informative estimate: it predicts a large but finite null space. The in-degree d_i (number of incoming edges) determines which neurons have degenerate incoming weights at $r = 45$. Neurons with $d_i \leq 45$ have enough independent activity dimensions to constrain all their incoming weights—they are “fully identifiable” and their weights can in principle be recovered from dynamics. Neurons with $d_i > 45$ have more incoming edges than independent activity dimensions, so $d_i - 45$ of their weights are free to vary without affecting the output—these are the null space dimensions where weight recovery is impossible:

In-degree range	Neurons	Null space contribution at rank 45
1–10	3,182	0 (fully identifiable)
11–20	3,917	0 (fully identifiable)
21–45	3,196	0 (fully identifiable)
46–100	2,843	bulk of null space
101–208	559	largest per-neuron null dimensions

High in-degrees arise because the 217 columns are laterally interconnected: a neuron integrating signals from its spatial neighborhood receives edges from the same cell type in many surrounding columns. Of the 13,741 neurons, 44 are source-only (in-degree zero, boundary neurons that project into the network but receive no edges). Summing over the remaining 13,697 postsynaptic neurons, **the SVD analysis predicts 115,223 null dimensions—26.5% of all 434,112 edge weights are unconstrained by the dynamics.**

2.3 Step 3: Within-type degeneracy (the dominant mechanism)

Step 2 showed *how large* the null space is (115,223 dimensions at rank 45) using only linear algebra and SVD. Step 3 explains *why* it is so large: the columnar organization of the connectome forces same-type neurons to have nearly identical activity, which is exactly what makes the columns of $\mathbf{H}_i$ linearly dependent and rank-deficient. Each of the 217 columns covers a different position in the visual field. For instance, L1 in column 5 sees the light at position (5, 3); L1 in column 6 sees position (6, 3). They receive different stimuli—but when a moving bar sweeps across the eye, column 6 sees what column 5 saw a moment earlier. Over thousands of frames with many stimuli, same-type neurons in different columns produce time-shifted copies of the same signal. Their activity is not identical, but highly correlated. Now consider a postsynaptic neuron i that receives input from k copies of the same cell type across k neighboring columns. Because these k neurons differ only by a small spatial shift, their activity vectors $\mathbf{h}_{j_1}(t), \dots, \mathbf{h}_{j_k}(t)$ are strongly correlated—like L1 in column 5 vs L1 in column 6, which differ by roughly one timestep. In the matrix $\mathbf{H}_i$, these k columns carry nearly the same information, so their rank is ≈ 1 instead of k , contributing $k - 1$ dimensions to $\ker(\mathbf{H}_i)$.

Why the sum must be constant. In the idealized case where the k columns are identical ($\mathbf{h}_{j_1} = \dots = \mathbf{h}_{j_k} = \mathbf{h}$), apply Step 2’s null space condition $\mathbf{H}_i \boldsymbol{\delta} = \mathbf{0}$. The contribution of these k edges is:

$$\delta_{j_1} \mathbf{h} + \delta_{j_2} \mathbf{h} + \dots + \delta_{j_k} \mathbf{h} = \left(\sum_{m=1}^k \delta_{j_m} \right) \mathbf{h} = \mathbf{0} \quad (9)$$

Since $\mathbf{h} \neq \mathbf{0}$, this requires $\sum_{m=1}^k \delta_{j_m} = 0$. So among k same-type edges, one constraint fixes the sum, leaving $k - 1$ free δ values—these are the null directions counted in Step 4.

2.4 Step 4: Counting the free parameters

For each postsynaptic neuron i and each presynaptic cell type α with $k_{i\alpha} > 1$ incoming edges, the null space gains $k_{i\alpha} - 1$ dimensions. Summing over all neurons and types:

$$\dim \ker(\mathbf{H}) = \sum_{i=1}^N \sum_{\alpha: k_{i\alpha} > 1} (k_{i\alpha} - 1) \quad (10)$$

For the flyvis connectome (13,741 neurons, 434,112 edges, 65 cell types):

Quantity	Value
Total edges E	434,112
Degenerate groups (dst, src_type with $k > 1$)	$\sim 121,000+$
Total null space dimension	$\sim 121,100$
Fraction of all edges in null space	$\sim 28\%$
Identifiable edge parameters	$\sim 313,000$

The null space has $\sim 121,100$ **dimensions**, meaning that many degrees of freedom in $\mathbf{W}$ are unconstrained by the dynamics. Within each group of k same-type edges targeting the same neuron, any perturbation satisfying the sum-zero constraint ($\sum \delta_{jm} = 0$, Step 3) leaves the dynamics unchanged. Weight can be freely *redistributed* among same-type inputs, but the total to each neuron from each type is fixed. Since each null direction admits a continuous scaling $\delta \rightarrow \lambda \delta$ for $\lambda \in \mathbb{R}$, the set of solutions is **uncountably infinite**: a $\sim 121,100$ -dimensional affine subspace of $\mathbb{R}^E$. This per-type structural count ($\sim 121,100$) agrees to within 5% with the global SVD bound from Step 2 (115,223). The global SVD captures all correlations (within- and cross-type) but is coarse; the per-type count precisely measures within-type redundancy but misses cross-type correlations. Their close agreement suggests that within-type degeneracy accounts for most of the null space.

3 Empirical Verification

The null space estimate in Step 2 relies on the effective rank at 99% variance—a conventional threshold, not an exact quantity. To test whether the predicted null directions are real (i.e., perturbations along them truly preserve dynamics), we verify empirically: **if** the null space analysis (Steps 2–4) is correct, **then** a perturbed matrix $\mathbf{W} + \delta$ with low connectivity R^2 (weights look very different) should still produce high rollout R^2 (dynamics look the same), as long as δ satisfies the sum-zero constraint ($\sum_{m=1}^k \delta_{jm} = 0$) within each same-type group. The flyvis connectome has 65 cell types, of which 52 are non-retina types that receive lateral inputs from same-type neurons across columns. Step 3 predicts two regimes depending on how correlated the same-type activity is:

- Types with **identical** same-type activity (within-column projections only) $\rightarrow$ **exact** null space, dynamics unchanged at any perturbation scale
- Types with **correlated but not identical** activity (cross-column projections) $\rightarrow$ **approximate** null space, small divergence growing with perturbation scale

Protocol. For each of the 52 non-retina types, we generated sum-preserving random perturbations at 15 scales (from $\lambda = 0.05$ to 8.0), producing 780 degenerate connectivity matrices. Each was rolled out through the full ODE from identical initial conditions and stimulus.

3.1 Key results

Metric	Value
Variants tested	780
Conn. R^2 range	0.28–1.00
Rollout R^2 range	0.80–1.00
Variants with conn. $R^2 < 0.9$ and rollout $R^2 > 0.99$	30 / 31 (97%)
Variants with conn. $R^2 < 0.99$ and rollout $R^2 > 0.999$	116 / 152 (76%)
Types with exact degeneracy (RMSE $\sim 10^{-7}$)	5 / 52
Types with approximate degeneracy	47 / 52

3.2 All 52 cell types at maximum perturbation ($\lambda = 8.0$)

Color code: **green** $R^2 > 0.9$, **orange** $R^2 > 0.5$, **red** $R^2 \leq 0.5$.

Cell type	Conn. R^2	Rollout R^2	Cell type	Conn. R^2	Rollout R^2
Tm20	0.282	0.997	Mi3	0.733	0.995
Tm5Y	0.566	0.998	Tm1	0.738	1.000
Mi1	0.617	0.998	Mi15	0.741	0.997
Lawf2	0.727	0.995	Tm4	0.753	0.986
Tm16	0.814	1.000	L3	0.871	0.999
Mi4	0.877	0.999	Lawf1	0.899	1.000
R1	0.915	0.975	R6	0.927	1.000
Tm5a	0.933	1.000	R2	0.938	0.998
TmY4	0.941	1.000	Mi13	0.942	1.000
R5	0.951	1.000	TmY5a	0.952	0.999
L2	0.965	1.000	T3	0.966	1.000
T4a	0.967	0.999	Mi14	0.967	1.000
T2a	0.969	1.000	T2	0.969	0.993
Tm2	0.969	1.000	TmY13	0.971	—
TmY14	0.972	1.000	TmY9	0.972	0.960
T4d	0.978	0.999	T4c	0.988	0.989
Tm3	0.990	1.000	T5c	0.990	1.000
T4b	0.990	1.000	T5d	0.992	1.000
Mi2	0.992	1.000	TmY3	0.993	1.000
TmY15	0.993	1.000	TmY10	0.994	1.000
Tm5b	0.995	1.000	Mi11	0.996	0.998
Tm30	0.997	1.000	R3	0.997	1.000
T5a	0.998	1.000	T5b	0.998	1.000
TmY18	0.998	1.000	L5	0.999	1.000*
L4	1.000	1.000*	R7	1.000	1.000*
R8	1.000	1.000*	Tm5c	1.000	1.000*

* Exact degeneracy: $\text{RMSE} \sim 10^{-7}$ (machine precision). — = diverged (TmY13 outlier). Sorted by connectivity R^2 (ascending = most degenerate first). The key pattern: the connectivity column shows **red** and **orange** (weights changed substantially) while the rollout column is entirely **green** (dynamics preserved). The 5 bold types at the bottom have exact degeneracy: their same-type columns in $\mathbf{H}_i$ are exactly identical, so $\mathbf{H}_i \boldsymbol{\delta} = \mathbf{0}$ holds exactly.

3.3 Two regimes

Exact degeneracy (5 types): L4, L5, R7, R8, Tm5c—within-column projections only, no spatial shift. The k same-type presynaptic neurons have truly identical activity. **Approximate degeneracy (47 types):** connectivity R^2 ranges from 0.28 (Tm20, 72% of weight variance changed) down to near 1.0, yet rollout R^2 stays above 0.96 in all cases. These results support the null space analysis: weights change measurably (most types conn. R^2 0.93–0.99, a few below 0.9) while dynamics remain nearly identical. The small residual divergence reflects the approximate nature of the null space— $\mathbf{H}_i \boldsymbol{\delta} \approx \mathbf{0}$, not exactly zero—due to the spatial shifts discussed in Step 3.

3.4 Per-scale summary

Each perturbation is $\mathbf{W} + \lambda \boldsymbol{\delta}$, where $\boldsymbol{\delta}$ is a random sum-preserving null direction and λ controls how far we move from the original weights. Larger λ means more different weights, lower connectivity R^2 .

λ	Conn. R^2 range	Mean rollout R^2
0.05	1.000–1.000	1.000000
0.50	0.997–1.000	0.999996
1.00	0.989–1.000	0.999984
2.00	0.955–1.000	0.999936
4.00	0.821–1.000	0.999722
8.00	0.282–1.000	$\sim 0.999^*$

* excluding one outlier (type 57, rollout $R^2 = 0.80$).

4 Discussion

4.1 The inverse problem is fundamentally ill-posed

Recovering the connectivity matrix $\mathbf{W}$ from dynamics alone is ill-posed, even with perfect noise-free observations, known τ_i , V_i^{rest} , known activation function, and known graph topology. The null space analysis (Steps 2–4) quantifies the ill-posedness: the weight space has $\sim 121,100$ **null dimensions** out of 434,112 edges ($\sim 28\%$), along which weights can be redistributed within same-type groups without affecting the dynamics. The empirical verification confirms this across all 52 non-retina cell types: at $\lambda = 8.0$, connectivity R^2 drops as low as 0.28 while rollout R^2 stays above 0.96. The set of valid solutions is **uncountably infinite**: the solution manifold is a $\sim 121,100$ -dimensional continuous subspace of $\mathbb{R}^{434,112}$. The ill-posedness is **structural**, arising from the columnar organization of the circuit: same-type neurons in different columns produce strongly correlated activity (Step 3), making their individual contributions to a shared target indistinguishable from dynamics alone.

4.2 Implications for GNN-based inference

The GNN approximates the dynamics as:

$$\frac{\widehat{dv}_i(t)}{dt} = f_\theta\left(v_i(t), \mathbf{a}_i, \sum_{j \in \mathcal{N}_i} \widehat{\mathbf{W}}_{ij} g_\phi(v_j(t), \mathbf{a}_j)^2, I_i(t)\right) \quad (11)$$

where $f_\theta = \text{MLP}_0$ and $g_\phi = \text{MLP}_1$ are three-layer perceptrons, $\widehat{\mathbf{W}}_{ij}$ are learnable edge weights, and $\mathbf{a}_i \in \mathbb{R}^2$ are learnable per-neuron embeddings. The training minimizes:

$$\mathcal{L}_{\text{pred}} = \sum_{i,t} \|\widehat{y}_i(t) - y_i(t)\|_2 \quad (12)$$

between simulator $y_i(t) = dv_i(t)/dt$ and GNN predictions $\widehat{y}_i(t) = \widehat{dv}_i(t)/dt$, augmented with regularization:

$$\begin{aligned} \mathcal{L} = & \|\widehat{\mathbf{y}} - \mathbf{y}\|_2 + \lambda_0 \|\theta\|_1 + \lambda_1 \|\phi\|_1 + \lambda_2 \|\widehat{\mathbf{W}}\|_1 \\ & + \gamma_0 \|\theta\|_2 + \gamma_1 \|\phi\|_2 \\ & + \mu_0 \left\| \text{ReLU}\left(-\frac{\partial g_\phi(v, \mathbf{a})}{\partial v}\right) \right\|_2 + \mu_1 \|g_\phi(v_\star, \mathbf{a}) - v_\star\|_2 \end{aligned} \quad (13)$$

where λ terms promote sparsity, γ terms stabilize learned functions, μ_0 enforces monotonicity of the edge message, and μ_1 normalizes with respect to reference voltage $v_\star$. When the GNN achieves $R_{\mathbf{W}}^2 = 0.997$ (as our LLM-optimized model does at $\sigma = 0.5$), it is recovering the $\sim 313,000$ identifiable parameters well while the $\sim 121,100$ null-space parameters are essentially free. The GNN’s implicit inductive bias (message-passing, weight sharing) acts as an implicit regularizer that selects a particular solution from the uncountable solution manifold.

4.3 Why noise helps

Adding process noise $\sigma\eta_i(t)$ to the simulation:

$$\tau_i \frac{dv_i(t)}{dt} = -v_i(t) + V_i^{\text{rest}} + \sum_{j \in \mathcal{N}_i} \mathbf{W}_{ij} \text{ReLU}(v_j(t)) + I_i(t) + \sigma\eta_i(t) \quad (14)$$

breaks the exact within-type degeneracy because:

1. Each neuron receives independent noise realizations, even if same-type
2. This injects independent variation into the columns of $\mathbf{H}_i$, breaking the within-type correlations from Step 3
3. The effective rank of $\mathbf{H}_i$ increases from 45 (noise-free) toward d_i
4. The null space $\dim \ker(\mathbf{H}_i) = d_i - \text{rank}(\mathbf{H}_i)$ shrinks, making more edge weights identifiable

Quantitative measurement. We measure how noise changes the effective rank of the population activity, and thus the null space dimension, using the same formula as Step 2: total null = $\sum_{i=1}^{N_{\text{post}}} \max(0, d_i - r)$. For each noise level ($\sigma = 0, 0.05, 0.5$), we:

1. Load the full voltage trace $\mathbf{V} \in \mathbb{R}^{64000 \times 13741}$ and subsample to 8,000 timesteps (every 8th frame)
2. Apply $h_j(t) = \text{ReLU}(v_j(t))$ to get the activity matrix $\mathbf{H} \in \mathbb{R}^{8000 \times 13741}$
3. Subsample to every 14th neuron $\rightarrow \mathbf{H}_{\text{sub}} \in \mathbb{R}^{8000 \times 982}$
4. Compute the **full** SVD of $\mathbf{H}_{\text{sub}}$ (all 982 singular values), so the total variance $\sum_k s_k^2$ is exact
5. Find the smallest r such that $\sum_{k=1}^r s_k^2 \geq 0.99 \cdot \sum_k s_k^2$
6. Plug r into the per-neuron null space formula above, using the in-degree distribution from the flyvis connectome (434,112 edges, max in-degree 208)

Results. The $\sigma = 0$ row reproduces Step 2’s estimate (rank 45, null 115,223) using the same SVD methodology. The noisy conditions use identical subsampling, so the ranks are directly comparable. GNN experimental results are from `flyvis_results.md` (10 seeds each):

Noise	Activity rank r	Null space dim	% identifiable	Observed $R_{\mathbf{W}}^2$
$\sigma = 0$	45	115,223	73.5%	0.899 ± 0.030
$\sigma = 0.05$	134	10,227	97.6%	0.962 ± 0.012
$\sigma = 0.5$	781	0	100.0%	0.997 ± 0.000

The trend validates the mechanism: noise increases the effective rank monotonically ($45 \rightarrow 134 \rightarrow 781$), shrinking the null space and making more weights identifiable. At $\sigma = 0.5$, the rank (781) exceeds the maximum in-degree (208), so every neuron’s incoming weights are fully identifiable—the null space vanishes entirely, consistent with $R_{\mathbf{W}}^2 = 0.997$. The two noisy conditions show good agreement between theoretical identifiability and observed recovery ($97.6\% \rightarrow 0.962$, $100\% \rightarrow 0.997$); the noise-free case ($R_{\mathbf{W}}^2 = 0.899$ vs. 73.5% identifiable) is higher than predicted, likely because GNN regularization biases null-space weights toward values that partially correlate with the ground truth. Alternatively, to break the strong within-type correlations experimentally, one could use optogenetics to inject independent perturbations into each neuron of a given type (e.g., stimulate each L1 with a different random light pattern), or inversely silence a subset of same-type neurons across columns (e.g., ablate L1 in columns 4 and 6 but not 5). Both approaches make the columns of $\mathbf{H}_i$ distinguishable, increasing $\text{rank}(\mathbf{H}_i)$ and shrinking the null space—the same mechanism as process noise, but achievable in real tissue.

5 Summary

Recovering the connectome from neural dynamics is fundamentally ill-posed: the columnar organization of the flyvis circuit forces same-type neurons across columns to produce strongly

correlated activity, making their individual synaptic contributions indistinguishable. Two independent estimates of the null space—a global SVD bound (115,223 dimensions) and a per-type structural count ($\sim 121,100$ dimensions)—agree to within 5%, confirming that within-type degeneracy is the dominant mechanism; the global bound captures all correlations (within- and cross-type) but is coarse, while the per-type count precisely measures within-type redundancy but misses cross-type correlations. This structural degeneracy affects 26.5% of all 434,112 edge weights, supported empirically by generating 780 perturbed connectivity matrices where weights change measurably (conn. R^2 ranging from 0.28 to 1.0, most between 0.93 and 0.99) while dynamics remain nearly identical (rollout $R^2 > 0.96$). Adding independent process noise breaks the within-type correlations and increases the effective activity rank from 45 (noise-free) to 781 ($\sigma = 0.5$), eliminating the null space entirely and enabling the GNN to achieve $R_{\mathbf{W}}^2 = 0.997$. The entire analysis rests on the rank-nullity theorem: a neuron’s incoming weights are recoverable when the activity rank exceeds its in-degree. Since the per-neuron rank cannot be measured directly, we upper-bound it by the global effective rank of the population activity (computed once by SVD); the only approximation is choosing the variance threshold that defines this effective rank. **The condition $r_{\text{global}} \geq d_i$ provides a principled, per-neuron criterion for when connectome recovery from dynamics is well-posed or ill-posed.**